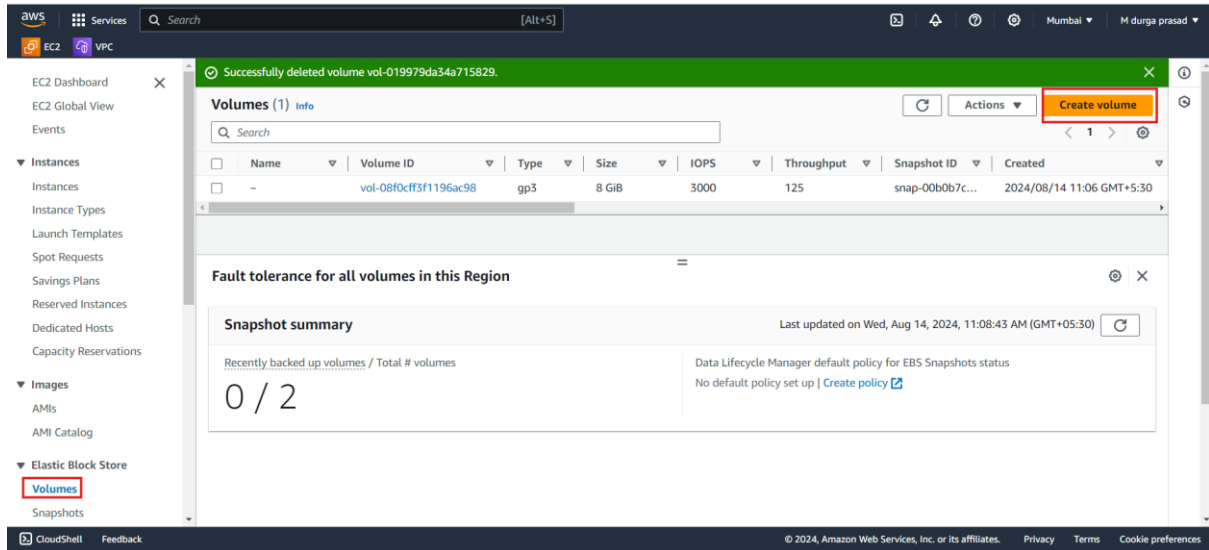


CONCAT LVM FILE SYSTEM IN AWS

First, we need to create volume as per requirement in volumes



Select volume type settings gp2, gp3, io1, io2, sc1, st1, standard

Volume settings

Volume type [Info](#)

General Purpose SSD (gp2)	▲
General Purpose SSD (gp3)	
General Purpose SSD (gp2)	✓
Provisioned IOPS SSD (io1)	
Provisioned IOPS SSD (io2)	
Cold HDD (sc1)	
Throughput Optimized HDD (st1)	
Magnetic (standard)	

[EC2](#) > [Volumes](#) > [Create volume](#)

Create volume [Info](#)

Create an Amazon EBS volume to attach to any EC2 instance in the same Availability Zone.

Volume settings

Volume type [Info](#)

General Purpose SSD (gp2)

Size (GiB) [Info](#)

100

Min: 1 GiB, Max: 16384 GiB. The value must be an integer.

IOPS [Info](#)

300 / 3000

Baseline of 3 IOPS per GiB with a minimum of 100 IOPS, burstable to 3000 IOPS.

Throughput (MiB/s) [Info](#)

Not applicable

Select available zone as per standard and also volume encrypt the as per requirement also take a snap shot for volume

Availability Zone [Info](#)

ap-south-1a

Snapshot ID - optional [Info](#)

Don't create volume from a snapshot

Encryption [Info](#)

Use Amazon EBS encryption as an encryption solution for your EBS resources associated with your EC2 instances.

☐ Encrypt this volume

Tags - optional [Info](#)

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

No tags associated with the resource.

Add tag

You can add 50 more tags.

Give a tap for easy to identifying as per standard

Snapshot summary [Info](#)

🕒 Click refresh to view backup information

The volume type that you select and the tags that you assign determine whether the volume will be backed up by any Data Lifecycle Manager policies.

Cancel

Create volume

Check the volumes tab the given volume create or not

✔ Successfully created volume vol-0c97aa4b09535a339.

Volumes (1/2) Info									
<div><div>🔍 Search</div></div>									
<input type="checkbox"/>	Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot ID	Created	
<input type="checkbox"/>	-	vol-08f0cff3f1196ac98	gp3	8 GiB	3000	125	snap-00b0b7c...	2024/08/14 11:06 GMT+5:30	
<input checked="" type="checkbox"/>	-	vol-0c97aa4b09535a339	gp2	1 GiB	100	-	-	2024/08/14 11:20 GMT+5:30	


check the volume tags as per requirement

Successfully created volume vol-0c97aa4b09535a339.

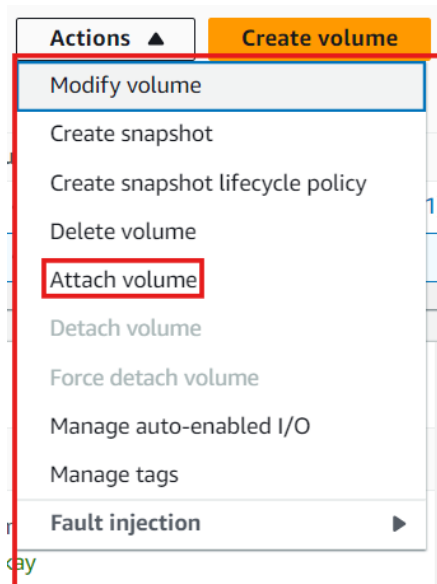
Volumes (1/2) Info						
Search						
Availability Zone	Volume state	Alarm status	Attached resources	Volume sta...	Encryption	KMS key ID
ap-south-1a	In-use	No alarms +	i-0155f92a582c12e50 (vara-sys-1): /dev/xvda (attached)	Okay	Not encrypted	-
ap-south-1a	Creating	No alarms +	-	Okay	Not encrypted	-

Check the volume state in available or pending

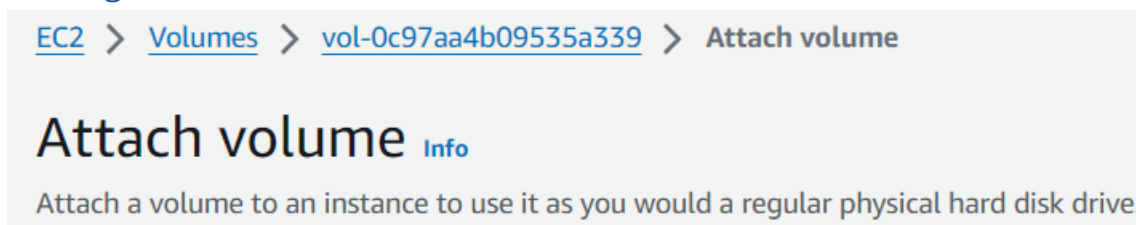
KMS key alias	Fast snaps...	Multi-Atta...
-	No	No
-	No	No



After that go to volume actions and attach the volume



Then go to the attach volume tab



Select the instance id which instance to attach as per running instance

The screenshot shows two parts of the AWS console. On the left, the 'Instance' dropdown menu is open, displaying a search bar and a list of instances. The instance 'i-0155f92a582c12e50 (vara-sys-1) (running)' is selected and highlighted with a red box. On the right, the 'Basic details' panel for a volume is shown, with a list of device paths. The path '/dev/sdb' is selected and highlighted with a red box.

Select the volume name as per standard & tab the attach volume

The screenshot shows the 'Attach volume' dialog in the AWS console. The 'Volume ID' is 'vol-0c97aa4b09535a339 (concat lv)'. The 'Availability Zone' is 'ap-south-1a'. The 'Instance' dropdown is set to 'i-0155f92a582c12e50'. The 'Device name' dropdown is set to '/dev/sdb'. A note states: 'Newer Linux kernels may rename your devices to /dev/xvdf through /dev/xvdp internally, even when the device name entered here (and shown in the details) is /dev/sdf through /dev/sdp.' At the bottom, there are 'Cancel' and 'Attach volume' buttons, with the 'Attach volume' button highlighted with a red box.

If you want to delete the volume then first need to detach or force detach

The screenshot shows the 'Force detach' dialog in the AWS console. The title is 'Force detach vol-0c97aa4b09535a339?'. The text explains that forced detachment can cause damage and should be used as a last resort. It asks, 'Are you sure that you want to force detach volume vol-0c97aa4b09535a339?'. Below this, it says 'To confirm detachment, type detach in this field.' A text input field contains the word 'detach' and is highlighted with a red box. At the bottom, there are 'Cancel' and 'Force detach' buttons, with the 'Force detach' button highlighted with a red box.

check the volume is detect or not in instance level in putty

by using the scan and lsblk

```
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# echo " - - - " > /sys/class/scsi_host/host0/s  
scan          sg_tablesize      subsystem/  
sg_prot tablesize state      supported mode  
[root@vara-sys-01 ~]# echo " - - - " > /sys/class/scsi_host/host0/scan  
[root@vara-sys-01 ~]# lsblk  
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS  
xvda         202:0    0   8G  0 disk  
├─xvda1      202:1    0   8G  0 part /  
├─xvda127    259:0    0    1M  0 part  
└─xvda128    259:1    0   10M  0 part /boot/efi  
xvdb         202:16   0    1G  0 disk  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]#
```

Then if want created into physical volume then command not because the lvm2 package not installed after that we need to install that required package by using yum module

```
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# pvcreate /dev/xvd  
xvda xvda1 xvda127 xvda128 xvdb  
[root@vara-sys-01 ~]# pvcreate /dev/xvdb  
-bash: pvcreate: command not found  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# yum clean all  
17 files removed  
[root@vara-sys-01 ~]# yum repolist  
repo id      repo name  
amazonlinux  Amazon Linux 2023 repository  
kernel-livepatch  Amazon Linux 2023 Kernel Livepatch repository  
[root@vara-sys-01 ~]# yum install lvm2  
Amazon Linux 2023 repository 50 MB/s | 26 MB  
[root@vara-sys-01 ~]#
```

#yum install lvm2

```
[root@vara-sys-01 ~]# yum repolist  
repo id      repo name  
amazonlinux  Amazon Linux 2023 repository  
kernel-livepatch  Amazon Linux 2023 Kernel Livepatch repository  
[root@vara-sys-01 ~]# yum install lvm2  
Amazon Linux 2023 repository 50 MB/s | 26 MB 00:00  
Amazon Linux 2023 Kernel Livepatch repository 58 kB/s | 11 kB 00:00  
Dependencies resolved.  
=====
```

Package	Architecture	Version	Repository	Size
Installing:				
lvm2	x86_64	2.03.16-1.amzn2023.0.5	amazonlinux	1.5 M
Installing dependencies:				
device-mapper-event	x86_64	1.02.185-1.amzn2023.0.5	amazonlinux	31 k
device-mapper-event-libs	x86_64	1.02.185-1.amzn2023.0.5	amazonlinux	30 k
device-mapper-persistent-data	x86_64	0.9.0-7.amzn2023.0.2	amazonlinux	781 k
lvm2-libs	x86_64	2.03.16-1.amzn2023.0.5	amazonlinux	985 k

```
=====
```

Transaction Summary

=====

Install 5 Packages

Total download size: 3.2 M
Installed size: 9.3 M
Is this ok [y/N]:

```

Is this ok [y/N]: y
Downloading Packages:
(1/5): device-mapper-event-libs-1.02.185-1.amzn2023.0.5.x86_64.rpm      482 kB/s | 30 kB    00:00
(2/5): device-mapper-event-1.02.185-1.amzn2023.0.5.x86_64.rpm        380 kB/s | 31 kB    00:00
(3/5): device-mapper-persistent-data-0.9.0-7.amzn2023.0.2.x86_64.rpm  7.1 MB/s | 781 kB   00:00
(4/5): lvm2-2.03.16-1.amzn2023.0.5.x86_64.rpm                       25 MB/s | 1.5 MB    00:00
(5/5): lvm2-libs-2.03.16-1.amzn2023.0.5.x86_64.rpm                   13 MB/s | 985 kB    00:00
-----
Total                                                                    15 MB/s | 3.2 MB    00:00
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing      :                                1/1
  Installing     : device-mapper-event-libs-1.02.185-1.amzn2023.0.5.x86_64 1/5
  Installing     : device-mapper-event-1.02.185-1.amzn2023.0.5.x86_64    2/5
  Running scriptlet: device-mapper-event-1.02.185-1.amzn2023.0.5.x86_64    2/5
Created symlink /etc/systemd/system/sockets.target.wants/dm-event.socket → /usr/lib/systemd/system/dm-event.socket.
  Installing     : lvm2-libs-2.03.16-1.amzn2023.0.5.x86_64                3/5
  Installing     : device-mapper-persistent-data-0.9.0-7.amzn2023.0.2.x86_64 4/5
  Installing     : lvm2-2.03.16-1.amzn2023.0.5.x86_64                    5/5
  Running scriptlet: lvm2-2.03.16-1.amzn2023.0.5.x86_64                    5/5
Created symlink /etc/systemd/system/sysinit.target.wants/lvm2-monitor.service → /usr/lib/systemd/system/lvm2-monitor.service.

```

Check the installation process

```

Created symlink /etc/systemd/system/sysinit.target.wants/lvm2-monitor.service → /usr/lib/systemd/system/lvm2-monitor.service.
Created symlink /etc/systemd/system/sysinit.target.wants/lvm2-lvmpolld.socket → /usr/lib/systemd/system/lvm2-lvmpolld.socket.

Verifying      : device-mapper-event-1.02.185-1.amzn2023.0.5.x86_64      1/5
Verifying      : device-mapper-event-libs-1.02.185-1.amzn2023.0.5.x86_64 2/5
Verifying      : device-mapper-persistent-data-0.9.0-7.amzn2023.0.2.x86_64 3/5
Verifying      : lvm2-2.03.16-1.amzn2023.0.5.x86_64                     4/5
Verifying      : lvm2-libs-2.03.16-1.amzn2023.0.5.x86_64                5/5

Installed:
device-mapper-event-1.02.185-1.amzn2023.0.5.x86_64
device-mapper-event-libs-1.02.185-1.amzn2023.0.5.x86_64
device-mapper-persistent-data-0.9.0-7.amzn2023.0.2.x86_64
lvm2-2.03.16-1.amzn2023.0.5.x86_64
lvm2-libs-2.03.16-1.amzn2023.0.5.x86_64

Complete!
[root@vara-sys-01 ~]#

```

After that installation create physical volume for that device check the #pvs, #pvdisplay, #pvs -v

```

[root@vara-sys-01 ~]#
[root@vara-sys-01 ~]# pvcreate /dev/xvdb
Physical volume "/dev/xvdb" successfully created.
[root@vara-sys-01 ~]#
[root@vara-sys-01 ~]# pvs
PV          VG Fmt Attr PSize PFree
/dev/sdb    lvm2 --- 1.00g 1.00g
[root@vara-sys-01 ~]# pvscan -v
PV /dev/sdb                                lvm2 [1.00 GiB]
Total: 1 [1.00 GiB] / in use: 0 [0 ] / in no VG: 1 [1.00 GiB]
[root@vara-sys-01 ~]#
[root@vara-sys-01 ~]# pvdisplay
"/dev/sdb" is a new physical volume of "1.00 GiB"
--- NEW Physical volume ---
PV Name           /dev/sdb
VG Name
PV Size           1.00 GiB
Allocatable       NO
PE Size           0
Total PE          0
Free PE           0
Allocated PE       0
PV UUID           9LD4Pi-Wv7v-Dk2C-Hd5y-Deyg-V9ie-6OHCcX
[root@vara-sys-01 ~]# pvs -v
PV          VG Fmt Attr PSize PFree DevSize PV UUID
/dev/sdb    lvm2 --- 1.00g 1.00g 1.00g 9LD4Pi-Wv7v-Dk2C-Hd5y-Deyg-V9ie-6OHCcX
[root@vara-sys-01 ~]#

```

Then create volume group from that volume & check the volume group is created or not by using # vgs, #vgs-v, #vgscan, #vgdisplay

```
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# vgcreate vg1 /dev/xvdb  
Volume group "vg1" successfully created  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# vgs  
VG   #PV #LV #SN Attr   VSize   VFree  
vg1   1   0   0 wz--n- 1020.00m 1020.00m  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# vgs -v  
VG   Attr   Ext   #PV #LV #SN VSize   VFree   VG   UUID                               VProfile  
vg1 wz--n- 4.00m   1   0   0 1020.00m 1020.00m aalbf5-6eEb-ZetD-dzjx-DZaV-mcAj-REBMSa  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# vgscan  
Found volume group "vg1" using metadata type lvm2  
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# vgscan -v  
Found volume group "vg1" using metadata type lvm2  
[root@vara-sys-01 ~]#
```

```
[root@vara-sys-01 ~]#  
[root@vara-sys-01 ~]# vgdisplay  
--- Volume group ---  
VG Name                vg1  
System ID  
Format                 lvm2  
Metadata Areas         1  
Metadata Sequence No   1  
VG Access               read/write  
VG Status               resizable  
MAX LV                 0  
Cur LV                 0  
Open LV                 0  
Max PV                 0  
Cur PV                 1  
Act PV                 1  
VG Size                 1020.00 MiB  
PE Size                 4.00 MiB  
Total PE                255  
Alloc PE / Size         0 / 0  
Free PE / Size          255 / 1020.00 MiB  
VG UUID                 aalbf5-6eEb-ZetD-dzjx-DZaV-mcAj-REBMSa  
[root@vara-sys-01 ~]#
```

After that create concat lv as per required size there are three manners

1. human readable # lvcreate -L +size(G,M) -n lvname vgname
2. by using physical extends # lvcreate -l PE's -n lvname vgname

```
[root@vara-sys-017M ~]#  
[root@vara-sys-017M ~]# lvcreate -l 100%FREE -n concatlv vg1  
Logical volume "concatlv" created.  
[root@vara-sys-017M ~]#  
[root@vara-sys-017M ~]# lvcreate -l 255 -n concatlv vg1  
Logical volume "concatlv" created.  
[root@vara-sys-017M ~]#
```


Then check the lv is created or not by using these commands

lvs, # lvscan, # lvs --segment, # lvdisplay

```
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# lvs
  LV VG Attr LSize Pool Origin Data# Meta# Move Cpy%Sync Convert
concatlv vg1 -wi-a----- 1020.00m
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# lvscan
[ACTIVE]                /dev/vg1/concatlv (1020.00 MiB) inherit
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# lvs -v
  LV VG #Seg Attr LSize  Maj Min RMaj RMin Pool Origin Data# Meta# Move Cpy%Sync Log Convert LV UUID LProfile
concatlv vg1 1 -wi-a----- 1020.00m -1 -1 253 0
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# lvdisplay
--- Logical volume ---
LV Path                /dev/vg1/concatlv
LV Name                 concatlv
VG Name                 vg1
LV UUID                 YeC62A-OTel-mlfZ-002j-idmw-ILNv-jUqCxG
LV Write Access         read/write
LV Creation host, time vara-sys-017M, 2024-08-14 09:08:25 +0000
LV Status               available
# open                  0
LV Size                 1020.00 MiB
Current LE              255
Segments                1
Allocation              inherit
Read ahead sectors      auto
 - currently set to     256
Block device            253:0

[root@vara-sys-017M ~]# lvs --segment
  LV VG Attr #Str Type SSize
concatlv vg1 -wi-a----- 1 linear 1020.00m
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]#
```

```
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# lvdisplay
--- Logical volume ---
LV Path                /dev/vg1/concatlv
LV Name                 concatlv
VG Name                 vg1
LV UUID                 YeC62A-OTel-mlfZ-002j-idmw-ILNv-jUqCxG
LV Write Access         read/write
LV Creation host, time vara-sys-017M, 2024-08-14 09:08:25 +0000
LV Status               available
# open                  0
LV Size                 1020.00 MiB
Current LE              255
Segments                1
Allocation              inherit
Read ahead sectors      auto
 - currently set to     256
Block device            253:0

[root@vara-sys-017M ~]# lvs --segment
  LV VG Attr #Str Type SSize
concatlv vg1 -wi-a----- 1 linear 1020.00m
```

If want to remove by using the # lvremove <device path>

```
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# lvremove /dev/vg1/concatlv
Do you really want to remove active logical volume vg1/concatlv? [y/n]: y
Logical volume "concatlv" successfully removed.
[root@vara-sys-017M ~]#
```


Then make file system in that logical volume and create mount point as per requirement

mkfs-t ext4 <concat lvm>

mkdir <mount point name> after that permanent in the file system tab file by using # vi /etc/fstab also check the # cat /etc/fstab

```
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# mkfs -t ext4 /dev/vg1/concatlv
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 261120 4k blocks and 65280 inodes
Filesystem UUID: 91752574-ec09-4c76-bf78-c709fc3824f0
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376

Allocating group tables: done
Writing inode tables: done
Creating journal (4096 blocks): done
Writing superblocks and filesystem accounting information: done

[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# mkdir /data
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# cat /etc/fstab
#
UUID=2277f5ea-eb9b-42da-a2e1-3b9cf1c1bca9 / xfs defaults,noatime 1 1
UUID=E239-DD44 /boot/efi vfat defaults,noatime,uid=0,gid=0,umask=0077,shortname=winnt,x-sys
temd.automount 0 2
/dev/vg1/concatlv /data ext4 defaults 0 0
[root@vara-sys-017M ~]#
```

Then mount the file system # mount <mount point name>

Check the file system mounted or not

By using # df-h or # df-h <mount point name>

```
[root@vara-sys-017M ~]# mount /data
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# df -h
Filesystem                Size      Used Avail Use% Mounted on
devtmpfs                  4.0M        0   4.0M   0% /dev
tmpfs                     475M        0   475M   0% /dev/shm
tmpfs                     190M    2.9M   188M   2% /run
/dev/xvda1                8.0G    1.6G   6.4G  20% /
tmpfs                     475M        0   475M   0% /tmp
/dev/xvda128              10M    1.3M   8.7M  13% /boot/efi
tmpfs                     95M        0    95M   0% /run/user/0
/dev/mapper/vg1-concatlv  986M     24K   919M   1% /data
[root@vara-sys-017M ~]#
[root@vara-sys-017M ~]# df -h /data
Filesystem                Size      Used Avail Use% Mounted on
/dev/mapper/vg1-concatlv  986M     24K   919M   1% /data
[root@vara-sys-017M ~]#
```